

OFFLOAD · ROUTING V2

WE BUILT THE RULER BEFORE THE WIN.

The v2 permit routing engine, and the board that grades it.
One route, one measure, one arm at a time. Never netted.
Where a ruler cannot speak, the board says so by name.

Offload · August 2026

WHAT WE ARE BUILDING, AND HOW WE KNOW

An engine that turns a printed oversize permit into a route a driver can follow, and a measurement system that grades every route it produces.

The engine

V2

Reads the permit PDF, extracts the printed route rows, and produces the driver line: waypoints, verbatim maneuvers, corridor holds. Three builds run side by side every time: v1 (ships today), v2_main (trunk), v2_pr (the change under test).

The corpus

401 ROUTES

Real permits from 32 states with real ground truth attached. Every run scores the whole corpus, not a sample.

The board

READ ONLY

Each run writes an immutable directory of results. The board is a web view over that archive. It never recomputes and it never nets scores into a single number.

The atom: one cell

A cell is one route, one measure, one arm. Ten routing measures, three arms, hundreds of routes. The board's whole design question is what to do with a cell nobody can grade yet, and its answer is to name it rather than drop it.

TWO QUESTIONS, ANSWERED ABOVE THE FOLD

Is the new engine better than the one that ships, and did the last run move anything.

VERSUS THE PREVIOUS GENERATION

v2_main meets or beats v1 on **1 of 8** graded measures (2 of 10 measures ungraded, so they are outside this count entirely rather than counted as losses)

v2_pr meets or beats v1 on **6 of 8** graded measures (2 of 10 measures ungraded, so they are outside this count entirely rather than counted as losses)

VERSUS THE LAST RUN

v2_main: **0 up** · **0 down** · 8 held · 2 not comparable (a measure with no scoreable set on one side is unknown, never a hold)


v2_pr: **0 up** · **0 down** · 8 held · 2 not comparable (a measure with no scoreable set on one side is unknown, never a hold)

differenced from the two runs' own published summaries, the same computation the compare view does; wave-c-comparands → wave-c-comparands-clean

v2_main 25.5% of the grid is ungraded on gradeable, made of 1,417 capture backlog (a capture run fixes these), 3,600 excluded by owner policy, and 1,424 blocked on truth or ruler work.

v2_pr 23.4% of the grid is ungraded on gradeable, made of 1,068 capture backlog (a capture run fixes these), 3,600 excluded by owner policy, and 2,412 blocked on truth or ruler work.

v2_main meets or beats v1 on 1 of 8 graded measures, v2_pr on 6 of 8. The 2 measures nobody can grade sit outside the count entirely, not scored as losses.



Where: the WHERE WE STAND band, top of the landing page.

A QUARTER OF THE GRID HAS NO VERDICT, AND IT SAYS SO

One alarm, split into three buckets you can act on separately. Only the first one is engineering backlog.

▲ **ROUTING LARGELY UNGRADED: 1417 of 5556 routing cells (25.5%) for v2_main have no graded comparison. Verdicts on this board are unreliable.**

1,417 **CAPTURE BACKLOG** **WORK** capture work EXISTS; a comparand re-capture fills these · 22% of the ungraded mass, 13.4% of all cells

3,600 **EXCLUDED BY OWNER POLICY** **OUT OF THE DENOMINATOR** WARM track. Out of every denominator by owner ruling 2026-08-06, kept and typed, never deleted · counted and typed, out of the alarm's a

1,424 **TRUTH / RULER LIMIT** **OUT OF THE DENOMINATOR** a TRUTH or RULER limitation, NOT backlog. Out of every denominator by owner approval 2026-08-06, kept and typed, never deleted · counted and typ

1,417 capture backlog (a re-capture run fills these). 3,600 excluded by owner policy. 1,424 blocked on truth or ruler work. Nothing is deleted: excluded cells stay on the board, typed and counted.

The exclusion is not the reason it looks good

The board prints what the number would read under the other two denominators: 40.7% under cold only, 60.9% under all rows. Both still trip the 25% threshold.

Typing the buckets moved the headline

It read 51.8% when every ungraded cell was one undifferentiated mass. It reads 23 to 25% now. That is a measurement fix, not an engine win, and the deck says so.

STOP 2 · THE SCORECARD

ONE ROW PER MEASURE. DENOMINATORS NEVER NETTED.

Each arm is scored over the set it could score, and the board refuses to average them into one figure.

ACCEPTANCE BAR

MEASURE	V1 (FLOOR)	V2_MAIN	V2_PR
backtrack	74/134	49/154	82/160
corridor	36/179	26/180	35/150
curfew	16/129	6/130	34/128
drivability	0/0	NO_RULER_FIXTURE	NO_RULER_FIXTURE
placement_binding	111/187	79/188	50/188
printed_text	14/170	14/163	23/160
served_pin	0/0	RULER_NOT_WIRED	RULER_NOT_WIRED
terminals	70/142	59/143	77/114
truth_bounds	35/142	34/143	51/114
verbatim	15/168	78/170	113/162

Where: the ACCEPTANCE BAR table on any run board.

+ detail

MEASURE	V1 (FLOOR)	V2_MAIN	V2_PR
backtrack	74/134	49/154	82/160
corridor	36/179	26/180	35/150
curfew	16/129	6/130	34/128
drivability	0/0	NO_RULER_FIXTURE	NO_RULER_FIXTURE
placement_binding	111/187	79/188	50/188
printed_text	14/170	14/163	23/160
served_pin	0/0	RULER_NOT_WIRED	RULER_NOT_WIRED
terminals	70/142	59/143	77/114
truth_bounds	35/142	34/143	51/114
verbatim	15/168	78/170	113/162

Passes over scoreable, per arm. verbatim reads 15/168, 78/170, 113/162.

verbatim is the measure that matters most to the product: it is the driver-facing printed text, and it is the one measure where both v2 arms clear v1 outright. **drivability** and **served_pin** read NO_RULER_FIXTURE and RULER_NOT_WIRED instead of 0, because a missing ruler is not a failing engine.

STOP 3 · THE MATRIX

THE WHOLE RUN AS ONE HONEST LEDGER

400 routes by 9 measures by 3 arms on the cold track, and a census that splits every cell by what happened to it.

Routes & measure matrix every column is the total count the column counts with relevant alarm always listed on bottom in this page or measure

MEASURE	ARM	PASS	FAIL	VALUE	REFUSED	GAP	N/A	UNDECIDABLE
backtrack	v1	94	39	0	0	213	54	0
backtrack	v2_main	83	70	0	0	212	35	0
backtrack	v2_pr	94	66	0	155	29	56	0
corridor	v1	0	130	0	0	222	47	1
corridor	v2_main	0	143	0	0	220	36	1
corridor	v2_pr	0	113	0	155	69	36	27
curfew	v1	0	101	28	0	239	32	0
curfew	v2_main	0	114	16	0	238	32	0
curfew	v2_pr	0	86	42	155	85	32	0
drivability	v1	0	0	0	0	400	0	0
drivability	v2_main	0	0	0	0	400	0	0
drivability	v2_pr	0	0	0	155	245	0	0
placement_binding	v1	0	17	169	0	214	0	0
placement_binding	v2_main	0	54	133	0	213	0	0
placement_binding	v2_pr	0	130	58	155	57	0	0
printed_text	v1	0	151	0	0	221	0	28
printed_text	v2_main	0	139	0	0	220	0	33
printed_text	v2_pr	0	136	0	155	68	0	33
served_pin	v1	0	0	0	0	400	0	0
served_pin	v2_main	0	0	0	0	400	0	0
served_pin	v2_pr	0	0	0	155	245	0	0
truth_bounds	v1	35	107	0	0	213	45	0
truth_bounds	v2_main	34	108	0	0	212	45	1
truth_bounds	v2_pr	50	63	0	155	66	44	22
verbatim	v1	29	139	0	0	232	0	0
verbatim	v2_main	142	28	0	0	230	0	0
verbatim	v2_pr	143	19	0	155	83	0	0

Where: the matrix view, one small square per cell, routes down the side and measure by arm across the top.

CELL CENSUS track cold · 400 routes x 9 measures x 3 arms

MEASURE	ARM	PASS	FAIL	VALUE	REFUSED	GAP	N/A	UNDECIDABLE
backtrack	v1	94	39	0	0	213	54	0
backtrack	v2_main	83	70	0	0	212	35	0
backtrack	v2_pr	94	66	0	155	29	56	0
corridor	v1	0	130	0	0	222	47	1
corridor	v2_main	0	143	0	0	220	36	1
corridor	v2_pr	0	113	0	155	69	36	27
curfew	v1	0	101	28	0	239	32	0
curfew	v2_main	0	114	16	0	238	32	0
curfew	v2_pr	0	86	42	155	85	32	0
drivability	v1	0	0	0	0	400	0	0
drivability	v2_main	0	0	0	0	400	0	0
drivability	v2_pr	0	0	0	155	245	0	0
placement_binding	v1	0	17	169	0	214	0	0
placement_binding	v2_main	0	54	133	0	213	0	0
placement_binding	v2_pr	0	130	58	155	57	0	0
printed_text	v1	0	151	0	0	221	0	28
printed_text	v2_main	0	139	0	0	220	0	33
printed_text	v2_pr	0	136	0	155	68	0	33
served_pin	v1	0	0	0	0	400	0	0
served_pin	v2_main	0	0	0	0	400	0	0
served_pin	v2_pr	0	0	0	155	245	0	0
truth_bounds	v1	35	107	0	0	213	45	0
truth_bounds	v2_main	34	108	0	0	212	45	1
truth_bounds	v2_pr	50	63	0	155	66	44	22
verbatim	v1	29	139	0	0	232	0	0
verbatim	v2_main	142	28	0	0	230	0	0
verbatim	v2_pr	143	19	0	155	83	0	0

CELL CENSUS: pass, fail, value, refused, gap, n/a, undecidable, per measure and arm.

Read the GAP column and the alarm stops being abstract: drivability is 400 gaps in every arm, served_pin the same. That is the 1,424 truth and ruler bucket, located.

A WAVE IS AN EXPERIMENT, AND IT CAN HONESTLY SHOW NOTHING MOVED

Runs are grouped into waves. Filter to one and every number recomputes over that wave, first run to last.

wave filter wave-c-comparands is on: 10 of 14 run(s) hidden. Every number on this page is computed over the 4 shown.

WHERE WE STAND the two questions every reader has, answered from the newest run's own published numbers. Everything below this band is the evidence for these two sentences

VERSUS THE PREVIOUS GENERATION

v2_main meets or beats v1 on **1 of 8** graded measures (2 of 10 measures ungraded, so they are outside this count entirely rather than counted as losses)

v2_pr meets or beats v1 on **6 of 8** graded measures (2 of 10 measures ungraded, so they are outside this count entirely rather than counted as losses)

ACROSS THIS WAVE WAVE-C-COMPARANDS, FIRST RUN TO LAST

v2_main: **0 up** · **0 down** · 8 held · 2 not comparable (a measure with no scoreable set on one side is unknown, never a hold)

v2_pr: **0 up** · **0 down** · 8 held · 2 not comparable (a measure with no scoreable set on one side is unknown, never a hold)

differenced from the two runs' own published summaries, the same computation the compare view does; the wave's first and last runs: wave-c-comparands → wave-c-comparands-clean

wave-c-comparands: 0 up, 0 down, 8 held, 2 not comparable, on both v2 arms. That wave was board and plumbing work, not engine work, so nothing should have moved. Movement here is the alarming outcome.

Routing board runs 14 completed runs · score moves only when a corpus run completes

all runs 14 board-rebuild 1 corpus-baseline-1 2 rulers-fold 1 substrate-freeze-era 2 wave-3-engine 3 wave-b-baseline 1 wave-c-comparands 4

The chip borders carry provenance: solid means the run producer stamped the label at render time, dotted means a human backfilled it afterwards. Same grouping, different epistemic weight.

THE BOARD WOULD RATHER REFUSE THAN HAND YOU A PLAUSIBLE WRONG NUMBER

Pick any two runs and the first thing it does is check whether they are comparable at all.


SUBSTRATE DIFFERS · these two runs are NOT comparable

different on: **panel model mtime, eval tree (the ruler)** · unknown on: osrm, osm. A delta measured across two substrates is a property of the substrate, not of the engine. Freeze one substrate and re-render both arms against it before reading any number below as a win or a regression.

FIELD	A (BASELINE)	B (CANDIDATE)	
cache root	/Users/nate/src/offload/v2-cache-qr-cold	/Users/nate/src/offload/v2-cache-qr-cold	SAME
freeze hash	a47ef410df47a28ba61f25ff08c7e6f7ef3ed1ae74922b6d0731a36d7dc268b3	a47ef410df47a28ba61f25ff08c7e6f7ef3ed1ae74922b6d0731a36d7dc268b3	SAME
panel model	/Users/nate/src/offload/panel-store-wavec/panel.js	/Users/nate/src/offload/panel-store-wavec/panel.js	SAME
panel model mtime	2026-08-07T02:22:45+00:00	2026-08-07T02:38:42+00:00	DIFFERS
arm v1 identity	bf405064d2d89e9aa69e	bf405064d2d89e9aa69e	SAME
arm v2_main identity	8b97dff33a@a8a688cdfb	8b97dff33a@a8a688cdfb	SAME
arm v2_pr identity	d8a57515e7@45ecd325f9	d8a57515e7@45ecd325f9	SAME
osrm (stamped by permit-service once task 19 lands; unknown until then)	unknown	unknown	UNKNOWN
osm (stamped by permit-service once task 19 lands; unknown until then)	unknown	unknown	UNKNOWN
truth fingerprint (the truth bytes each run scored against; a difference makes the delta truth-attributable)	ea453d66a5feeb2315fdacc47cc3c5d69db4f45320d4ff4f6153b0e28adb158c	ea453d66a5feeb2315fdacc47cc3c5d69db4f45320d4ff4f6153b0e28adb158c	SAME
eval tree (the ruler) (git tree of eval/; a difference makes the delta scorer-attributable)	642d239b35b0687bffd9fe3013fdbc155200a64	9b574a251c358ea699656a3b5ee144a4c841eb3	DIFFERS

identity is CONSUMED here, never recomputed: each field is what that run recorded at render time. The substrate VERDICT belongs to eval/verbatim_gate.py substrate_decision in permit-service; this board never forms a second one.

SUBSTRATE DIFFERS, opened in red, with the receipts underneath: cache root same, freeze hash same, all three arm identities same, eval tree differs.



Where: the compare view header. Further down, the movers ranking collapses to NOT ATTRIBUTABLE instead of showing a leaderboard.

EVERY VERDICT, WITH THE TYPED REASON PRINTED INLINE

Route 1528, Arkansas. The two v2 arms disagree, so the page has to explain itself rather than assert.

ROUTE VERDICT graded cells only · badge requires the placement + terminals + truth_bounds core graded on both arms · direction = graded W vs L, TIE only when W = L, both-fail never credits


v2_main · LOSS v2_pr · WIN

PER-MEASURE OUTCOME, COLD TRACK MEETS / MISSES / typed refusal, straight off absolute.outcomes · the typed reason is printed inline, not only on hover

MEASURE	V1	V2_MAIN	V2_PR	TYPED REASON (V2_PR)
backtrack	MEETS 1✓	MEETS 1✓	MEETS 1✓	
corridor	MISSES 3✓ 1x	MISSES 0✓ 4x	MEETS 1✓	
curfew	ZERO_DENOMINATOR	ZERO_DENOMINATOR	ZERO_DENOMINATOR	ZERO_DENOMINATOR this permit prints 0 curfew span(s) and 0 of them are single-window expressible, so structured recall has no denominator here
drivability	NO_RULER_FIXTURE	NO_RULER_FIXTURE	NO_RULER_FIXTURE	NO_RULER_FIXTURE no drivability report for arm v2_pr at /Users/nate/src/offload/_wt/wt-board-build/eval/artifacts/team3/reports/drivability_1528.json; produce one with 'python -m eval.served_capture_run' then 'python -m eval.team3_drivability' (v1 arm: 'eval.v1_floor_harvest --emit-v1-envelopes' + TEAM3_SCRATCH=eval/artifacts/team3_v1)
placement_binding	MEETS 1✓	MISSES 0✓ 1x	MEETS 1✓	
printed_text	MEETS 3✓	MISSES 0✓ 3x	MISSES 2✓ 1x	FAILED count
served_pin	RULER_NOT_WIRED	RULER_NOT_WIRED	RULER_NOT_WIRED	RULER_NOT_WIRED no OSM_DB.* in this environment, so the printed refs cannot be resolved against live OSM at score time. The graded surface for this measure is eval.served_c2_gate (committed baseline over the HEAD served capture, local lab); this cell fills when the panel runs with the lab env exported.
terminals	MEETS 2✓	MEETS 2✓	MEETS 2✓	
truth_bounds	MEETS 5✓	MEETS 5✓	MEETS 2✓	
verbatim	MISSES 0✓ 1x	MISSES 0✓ 1x	MISSES 0✓ 1x	FAILED verbatim

MEETS and MISSES per arm, and the right column names why: ZERO_DENOMINATOR, NO_RULER_FIXTURE with the exact missing artefact path and the two commands that would produce it.

This is refusal as product. The board does not just say a cell is ungraded, it says who can close it and how.



Where: the route drill-down.

AN ALARM NOBODY CAN ACT ON IS JUST NOISE

The 1,424 truth and ruler cells, sorted into classes and priced by artefact, with the team that fetches each one named.

THE CLASSES the producer's own cell kind plus a three-entry code override; no message text is parsed

ONBOARDABLE 1,499 cell(s) truth is obtainable; onboarding it grades these cells

STRUCTURALLY INAPPLICABLE 476 cell(s) 0/0 by a property of this route; permanently ungradeable, NOT workload

DELIVERY DEPENDENT 435 cell(s) the ruler ran and could not decide on what was delivered

RULER DEFECT 2 cell(s) the ruler raised; a bug, not missing truth

WHAT IS MISSING by artefact, with the team that fetches it; a pipeline artefact is not expert time

TRUTH ARTEFACT	OWNER	CELLS	ROUTES	MEASURES IT UNLOCKS	WHERE IT LIVES
expert truth REPAIR	expert	308	44	truth_bounds	expert truth exists for this route and fails its own admissibility bounds; it needs
expert route truth	expert	259	37	truth_bounds	the hand-corrected expert route file for this route (route_<id>.json in the truth co
corridor ruler fixture	expert	234	39	corridor	eval/team1_fixtures/corridor/route_<id>.json, the C1-C6 placement ruler
ruler service wiring	ops	216	216	served_pin	the ruler is not connected to the service it reads (typically no OSM_DB_* in this er
drivability report	pipeline	216	216	drivability	eval/artifacts/team3/reports/drivability_<id>.json, produced by eval.served_capture_
printed-maneuver fixture	harvest	156	39	printed_text	eval/pt_fixtures/printed_<id>.json, harvested with eval.pt_harvest
printed-curfew fixture	harvest	56	56	curfew	eval/cf_fixtures/printed/<id>.json, harvested with eval.cf_printed_curfews
persisted page blocks	harvest	54	54	verbatim	the Textract blocks for this permit; without them the printed page cannot be reconst

1,499 onboardable (the only class that is work). 476 structurally inapplicable (0/0 by a property of the route, permanently ungradeable, never workload). 435 delivery dependent. 2 ruler defect. Below, the same mass priced by artefact and owner.

Of everything the alarm counts, **801 cells across 81 routes** are genuinely expert time. The rest is ops wiring, pipeline output, or harvest.

HOW MUCH OF THIS IS REAL

The honest answer to "what are you actually grading against", stated as an identity rather than a claim.

```
COMPOSITION  n_expert_only + n_qr_only + n_both_agree + n_both_conflict = n_total by construction

TOTAL                EXPERT ONLY                QR ONLY                BOTH, AGREE
401                  190                        211                    0
HUMAN APPROVED      QR SOURCED                QR PLAUSIBILITY FLAG    STATES
191                  211                        60                    32

scored by tier: {"expert_approved":190,"qr_official":210,"qr_state":1}
by state: {"Alabama":14,"Arizona":1,"Arkansas":29,"California":8,"Colorado":20,"Florida":1,"Georgia":40,"Idaho":1,"Illinois":5,"Indiana":3,"Iowa":6,"Kansas":8,"Kentucky":4,"Louisiana":2,"Maryland":4,"Minnesota":2,"Mississippi":10,"Ohio":2,"Oklahoma":11,"Pennsylvania":3,"South Carolina":1,"South Dakota":2,"Tennessee":23,"Texas":141,"Virginia":3,"West Virginia":5,"Wisconsin":5}

QR FLAGGED  60 routes · carries QR_PLAUSIBILITY_FLAG: unverifiable, and IN the cohort
1099 · 1118 · 1146 · 1147 · 1149 · 1365 · 1451 · 1452 · 1529 · 1593 · 1595 · 1731 · 1796 · 1997 · 2000 · 2066 · 2228 · 2252 · 2277 · 2279 · 2291 · 2296 · 2394 · 2452 · 2456 · 2648 · 2676 · 2686 · 2725 · 2814 · 2869 · 3053 · 3515 · 3803 · 3845 · 3944 · 3950 · 3986 · 3996 · 4004 · 4326 · 4433 · 4481 · 4583 · 4601 · 4637 · 4664

expert-only and qr-only are published as COUNTS, not as route lists, so they cannot be filtered without inventing membership. That absence is a fact about the model, not about the cohort.
```

401 routes: 190 expert only (a human read the permit and hand-corrected the route), 211 QR only (the issuing state's own route service), 0 both agree, 0 both conflict. 32 states, led by Georgia 40, Arkansas 29, Colorado 20.

The gap this exposes

0 both-agree means no route currently carries both kinds of truth, so expert truth and state QR truth have never been cross-checked against each other. That is a known hole, stated on the page.

Flagged and kept

60 QR records carry a plausibility flag: unverifiable, but kept in the cohort and listed by id rather than quietly dropped.

COLD ONLY: THE ENGINE NEVER SEES THE ANSWER

Cold means the curated dictionary and corridor services are unset. It answers the load-bearing question: does the engine work from nothing on a permit nobody has ever curated.

OWNER RULING 2026-08-06: the warm track is not a track we run, and it is out of every ungraded

denominator. In his words: "warm track? We will never run a warm track, this sounds like it's polluting our results.

Maybe later we will want to include warm but not today." Read it against the definitions above: the warm track wires the dictionary and corridor services INTO the engine, which makes expert data an **engine INPUT** rather than a ruler.

The cold track is the one we grade, and it is the one that carries the claim that the engine works from nothing.

	COLD TRACK	WARM TRACK
expert data (dictionary / corridor)	not wired. The engine has never seen it	wired INTO the engine as an input
what a score there means	the engine resolved the permit on its own	the engine was handed part of the answer
measured today?	YES. This is the campaign	NO, by owner ruling 2026-08-06. Captured as typed absences, kept on the board, out of every denominator
on the board	every rate, every banner, every alarm	its own banner line and its own census line, marked OUT OF THE DENOMINATOR


The glossary carries the ruling that shapes every number on the board, and is explicit that this is a denominator decision and not a deletion.

Why warm cells are refused rather than absent

On the warm track, corridor, served_pin and drivability would be scored against the very tables that were fed to the engine. That is circular by construction, so those cells are typed REFUSED. Expert data is a first class input to the product; it is never a ruler.

WHAT THE DRIVER ACTUALLY GETS

The board has no map by design. Route geometry is rendered separately, from the same banked capture the board scored.



Texas permit 260714107904: 105 pins, 111 driver rows, 1213.5 km, with the stretches that leave the state's own QR route overlaid in red.

REFUSALS BY TIER (placement_binding_audit, 105 driver-visible pins)

engine_solution	18 / 105 pins
geometric_fallback	40 / 105 pins
printed_feature	21 / 105 pins
printed_mileage	26 / 105 pins
refused	3 / 105 pins
repaired_geometry	0 / 105 pins

Every pin carries how it was placed, and 3 of 105 are an outright refusal rather than a guess.

Rendered, not asserted

Basemap, pins, verdict chip and the machine facts footer all come off one immutable capture. No engine ran to make this page.

THE DISCIPLINE BEHIND EVERY NUMBER IN THIS DECK

Four rules. Each one exists because breaking it cost us a day.

Freeze the substrate before an A/B

One live expert fanout at the target commit, immutable from then on. The warm cache drifts and is mutated by capture, so a fresh arm against a stale arm measures the cache, not the engine. We run the same code control first, every time.

Measure at \$0 by replay

Fingerprint keyed extraction replays off the frozen capture. A 6 route probe landed 5 routes in about 70 seconds with full routing callbacks and no billable calls, against 0 of 400 in 25 minutes before that fix. The spend guard is armed on every run.

Verify adversarially before folding

Every candidate is graded by an independent lane that did not build it, on the frozen substrate, before it folds into trunk. An agent report is a claim until the gate reproduces it.

Bank the run, keep the provenance

Every run writes an immutable directory: arm identities, cache root, freeze hash, truth fingerprint, eval tree. That is what lets the compare view refuse. Nothing on the board is recomputed at read time.

WHERE IT STANDS

WHAT IS LANDED, WHAT IS IN FLIGHT

Honest status. The engine numbers on this board are current; the next set is still being scored.

PR #210

CARRIES IT

The v2 engine plus the whole measurement layer: the board, the cell census, the truth queue, the compare gate, the wave grouping and the typed refusals. One review surface, one story.

Corpus re - measure

CAPTURE DONE

246 routes captured at \$0 against a read only frozen substrate: 214 decision relevant plus a 32 route stratified sample of the inert class, so inert is confirmed rather than assumed. Per - measure scoring is in measurement.

Engine candidates

IN MEASUREMENT

Improvement candidates exist and are being graded by an independent lane on the frozen substrate. No number goes in a deck until that lane reports.

Known and being fixed in parallel

The matrix grid paints one colour for every state (a CSS specificity bug, census numbers unaffected). No run in the archive carries a blessed baseline yet, so the per - measure delta column has never been exercised. OSRM and OSM substrate identity reads unknown until the stamping task lands, which is what forces the compare view to refuse.

THE POINT

BUILD THE RULER FIRST

Honest beats flattering

Per measure, per arm, never netted. Two measures with no ruler are held out of the count rather than scored as zeros, and the board prints what the alarm would read under every other denominator.

A refusal is a product

"Could not resolve, and why" reaches the page as typed data with the artefact path and the command that fixes it. That is the provable half of the north star.

An alarm becomes a work list

1,424 ungraded cells resolve to 801 cells across 81 routes of real expert time, priced by artefact, with an owner on each row.

We can now tell the truth about whether a change to the engine helped, per measure, with receipts. That is the thing that was missing.

Every screenshot in this deck is the live board over the 14 run local archive, plus one route render off an immutable capture. Nothing here is a mockup.